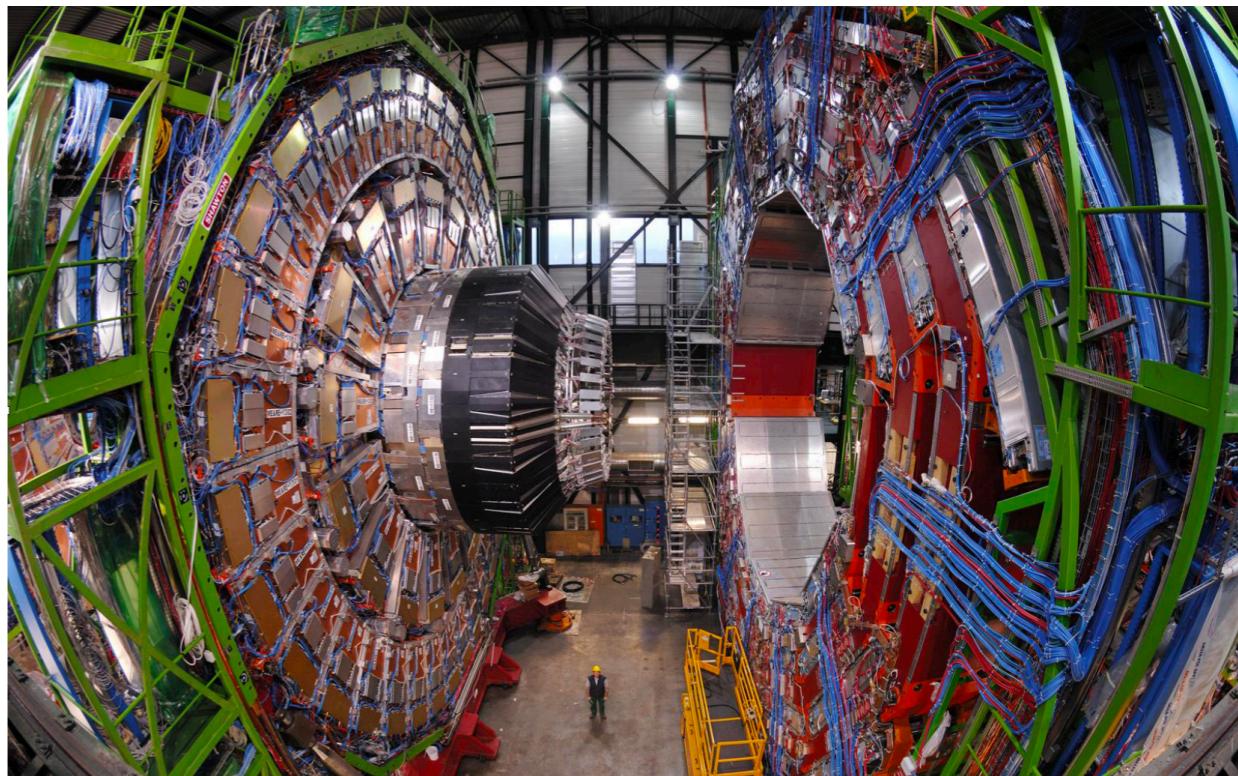
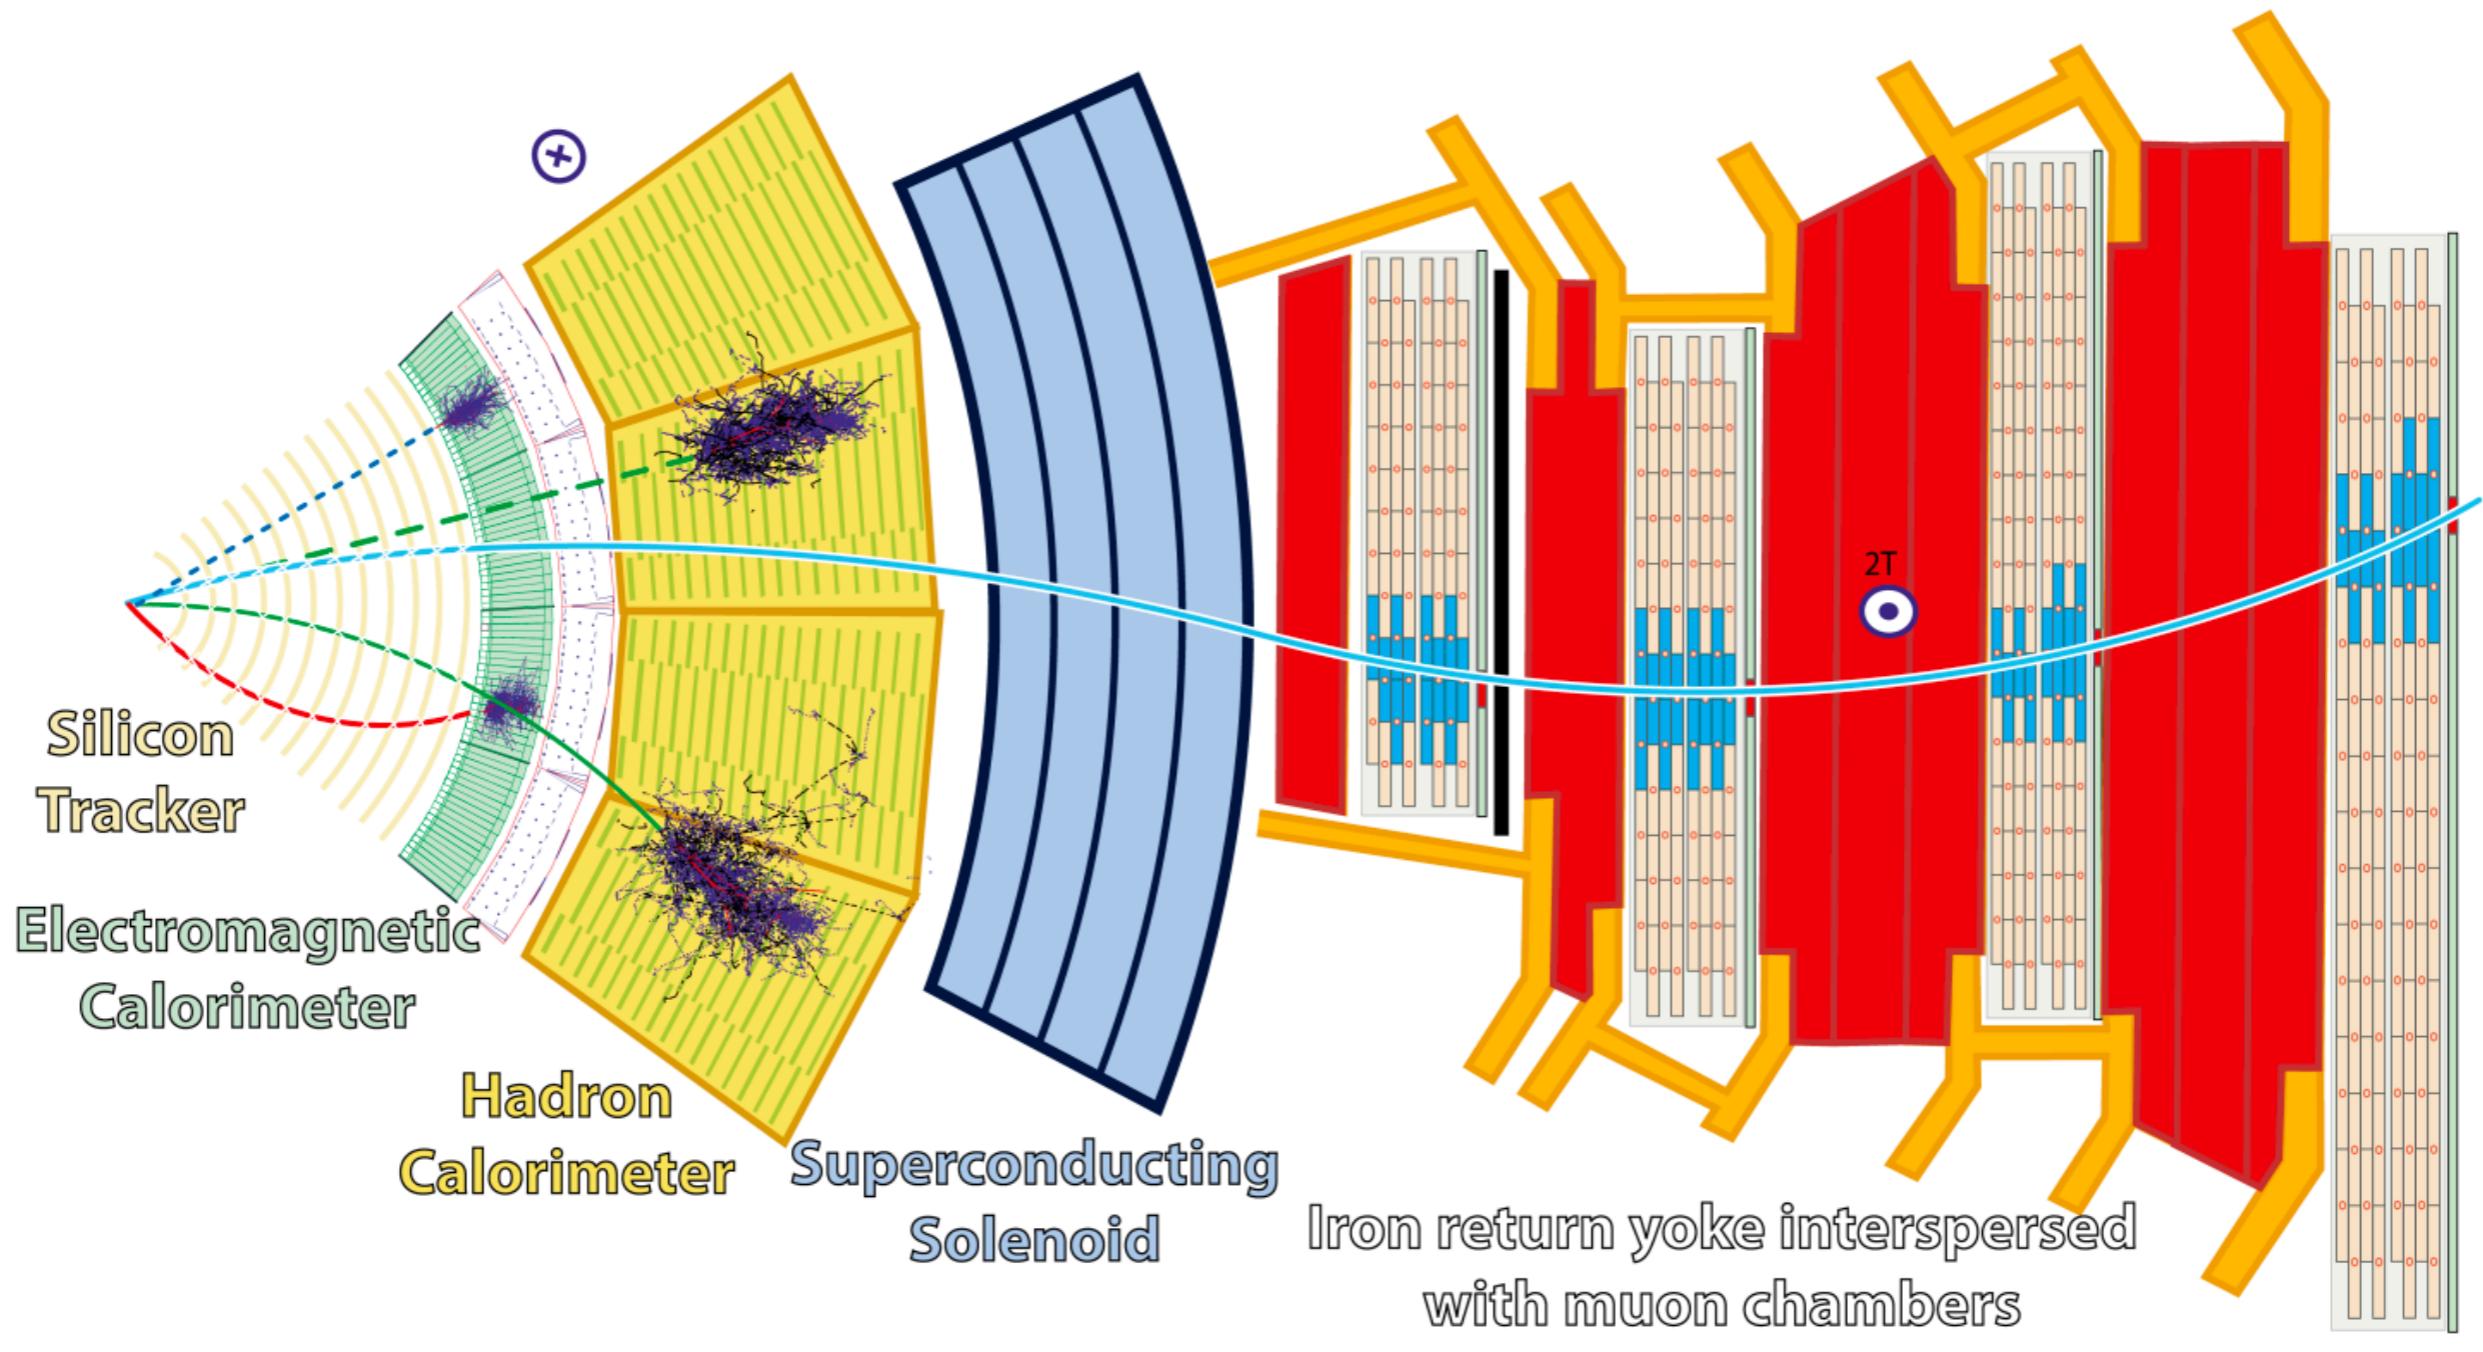


Large Hadron Collider

- Large Hadron Collider (proton - proton)
- Dedicated detectors to record data from collisions
 - Electromagnetic calorimeter (ECAL) is designed to measure energy of EM particles, $e + \gamma$ (**EG**)
- Many overlapping collisions in addition to the primary one, called **pileup (PU)**
- Want to design algorithm to distinguish primary **EG** energy deposits from **PU**



CMS Detector



— Muon

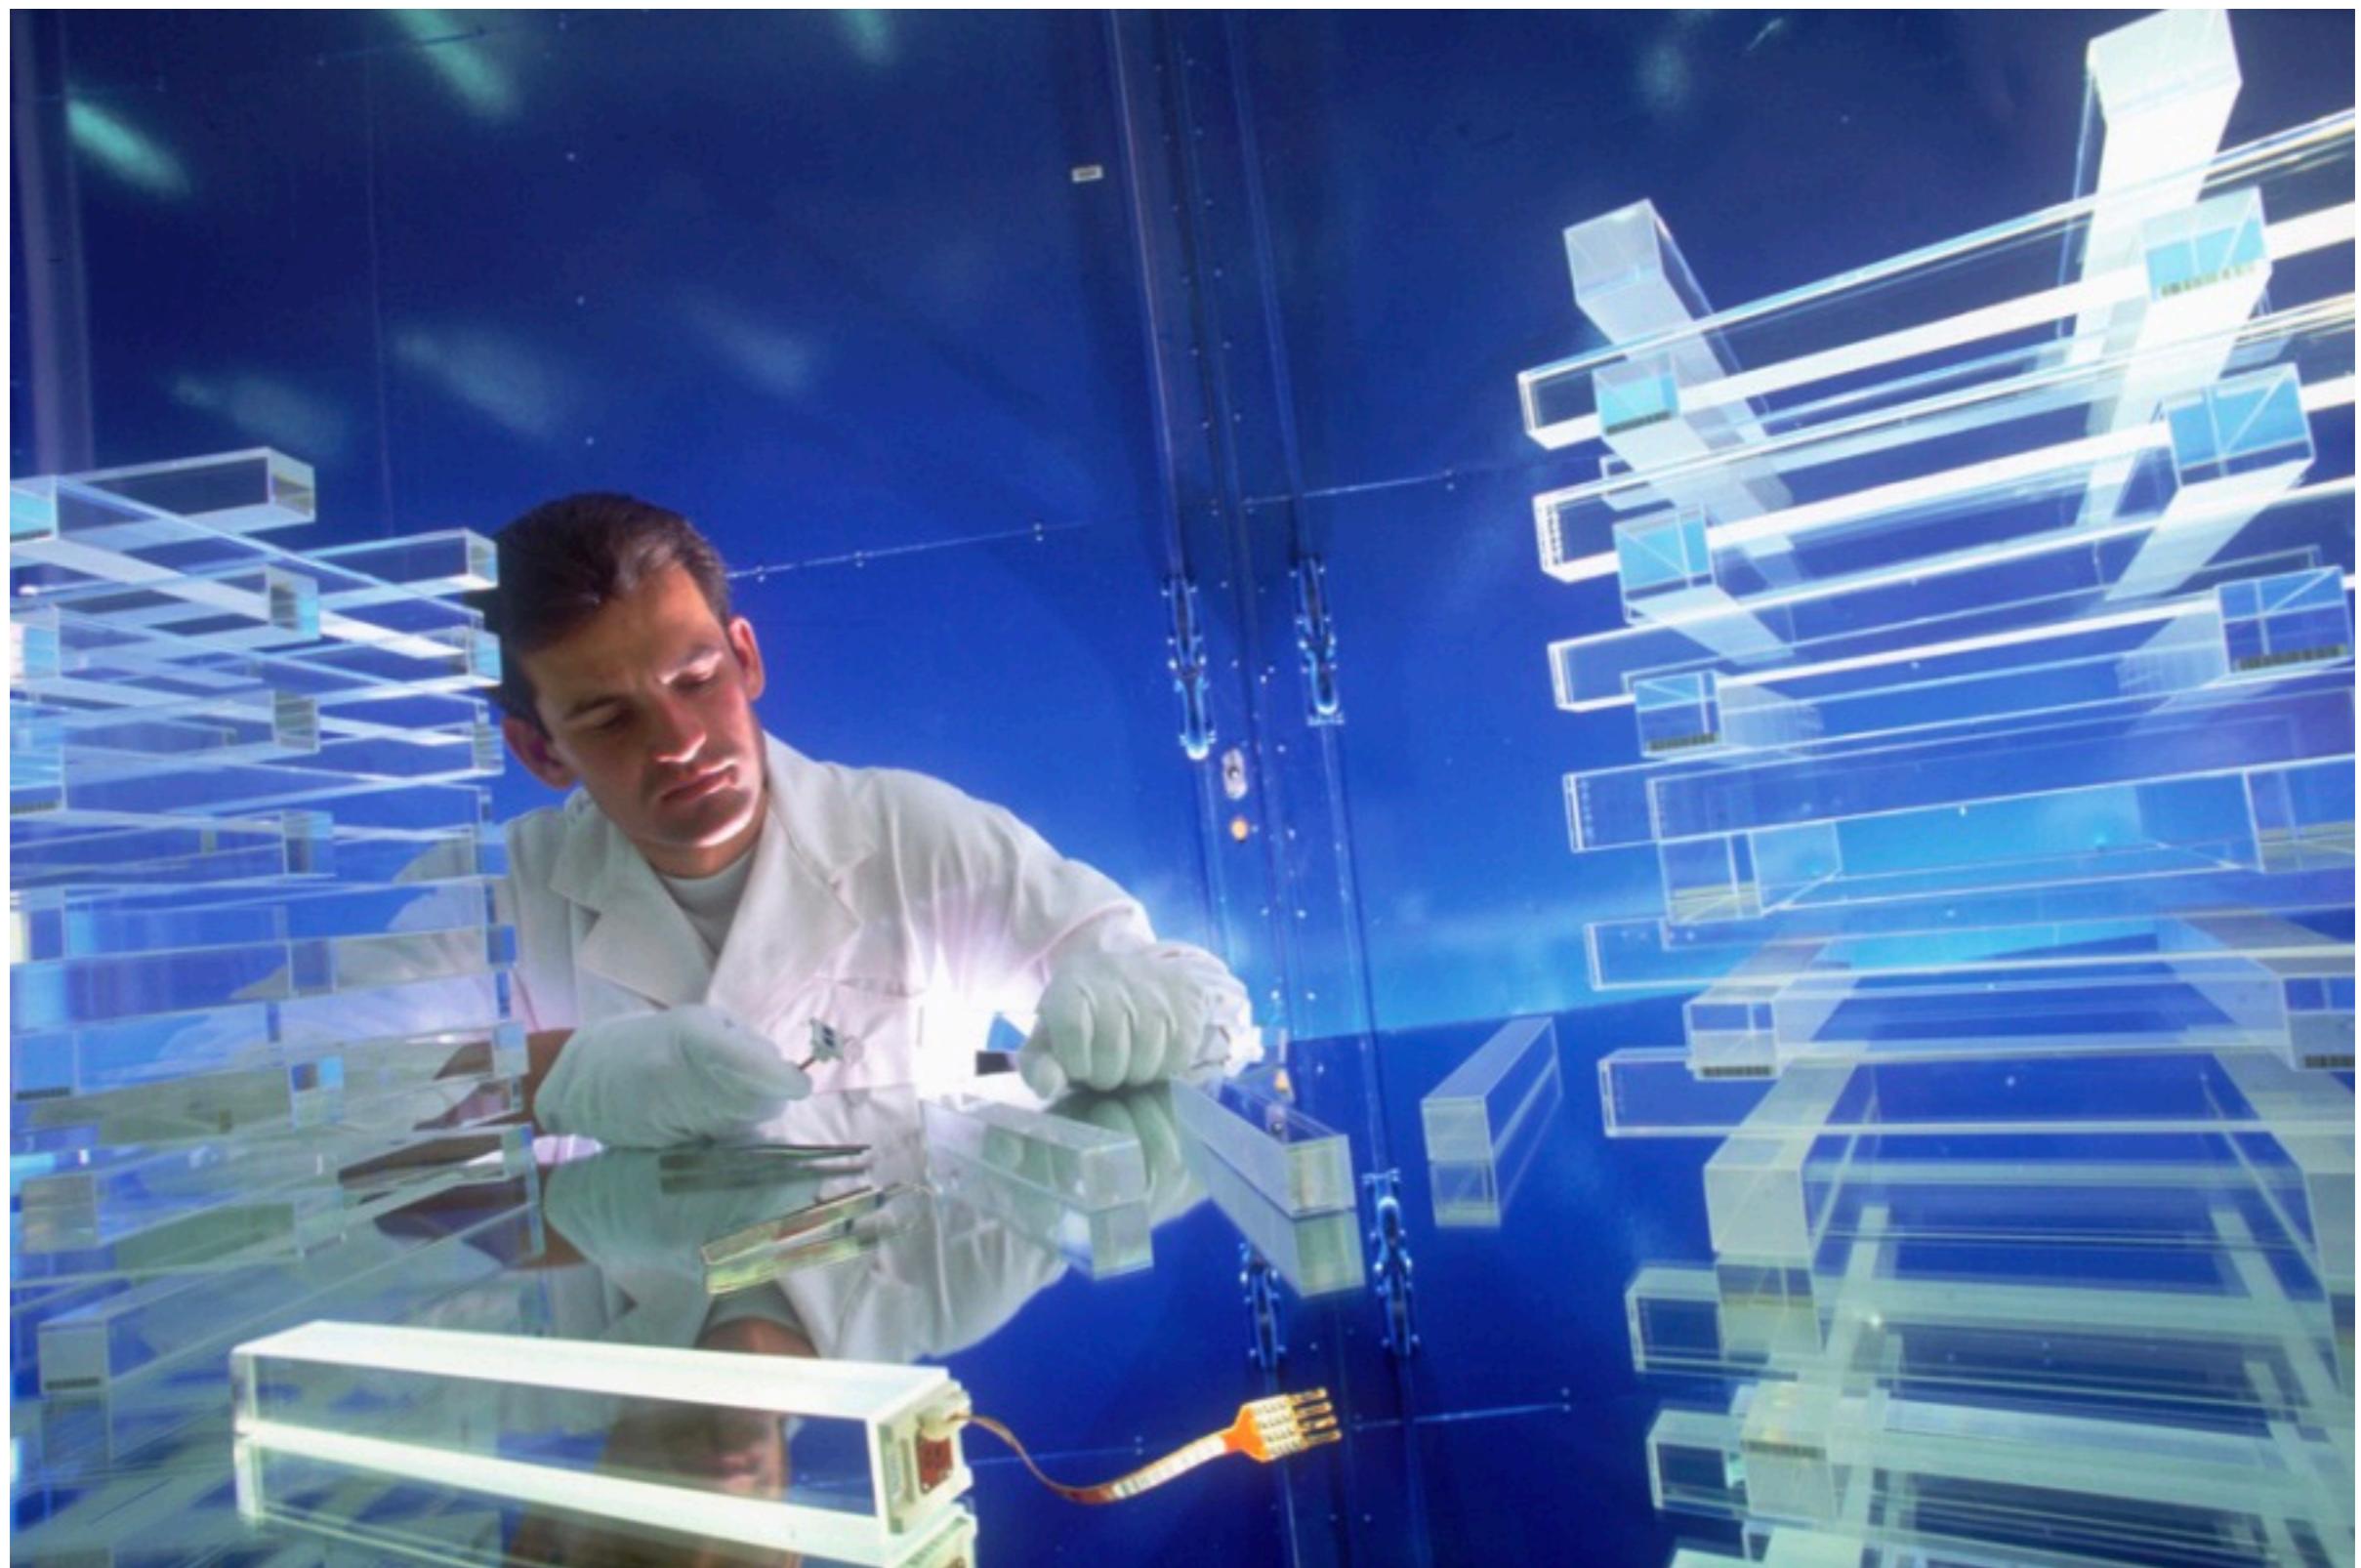
— Electron

— Charged hadron (e.g. pion)

— Neutral hadron (e.g. neutron)

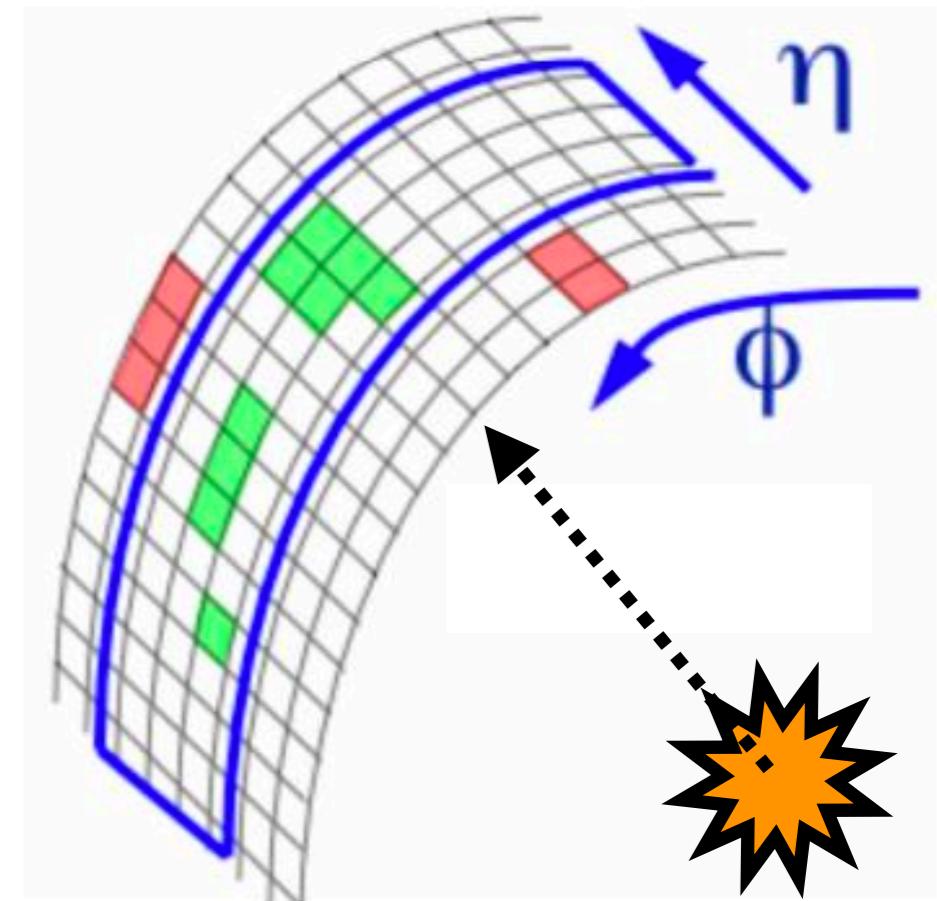
----- Photon

CMS ECAL

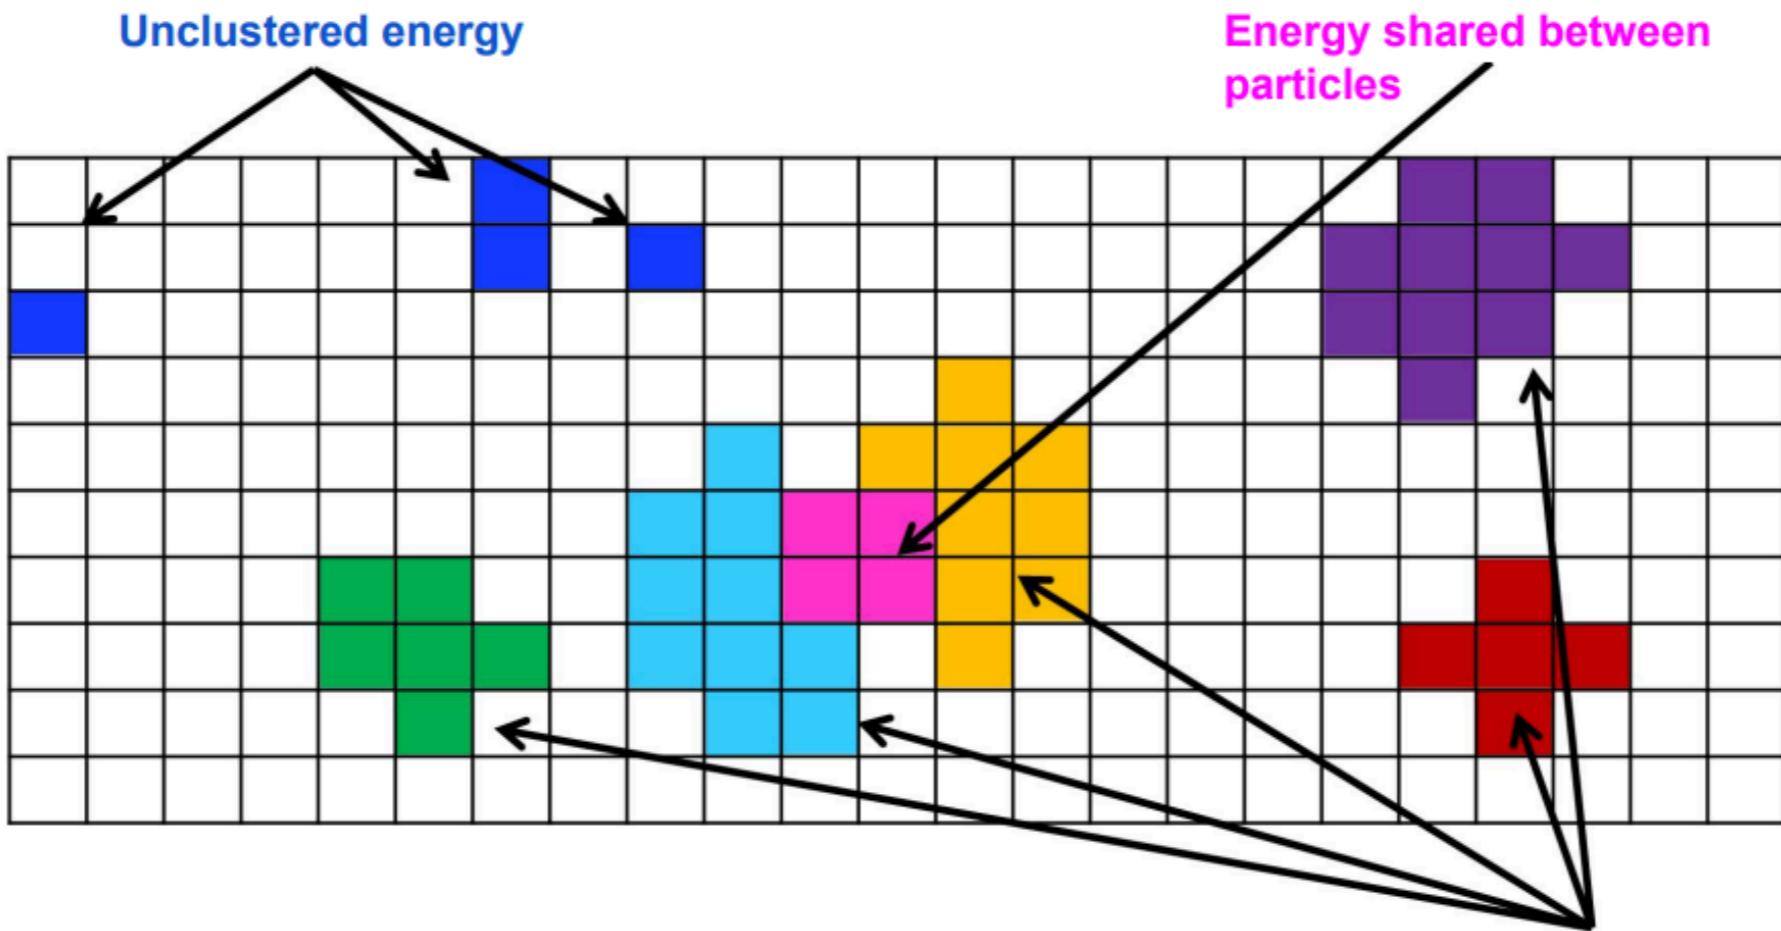


Dataset Explanation

- Real photon and electrons will deposit energy in certain patterns
 - PU is either random or from other sources, should not match
- We describe energy deposits using “shower shape variables”
 - i.e. How much energy is in certain regions around the center? How correlated are deposits in η and ϕ ? How many crystals in particular directions?

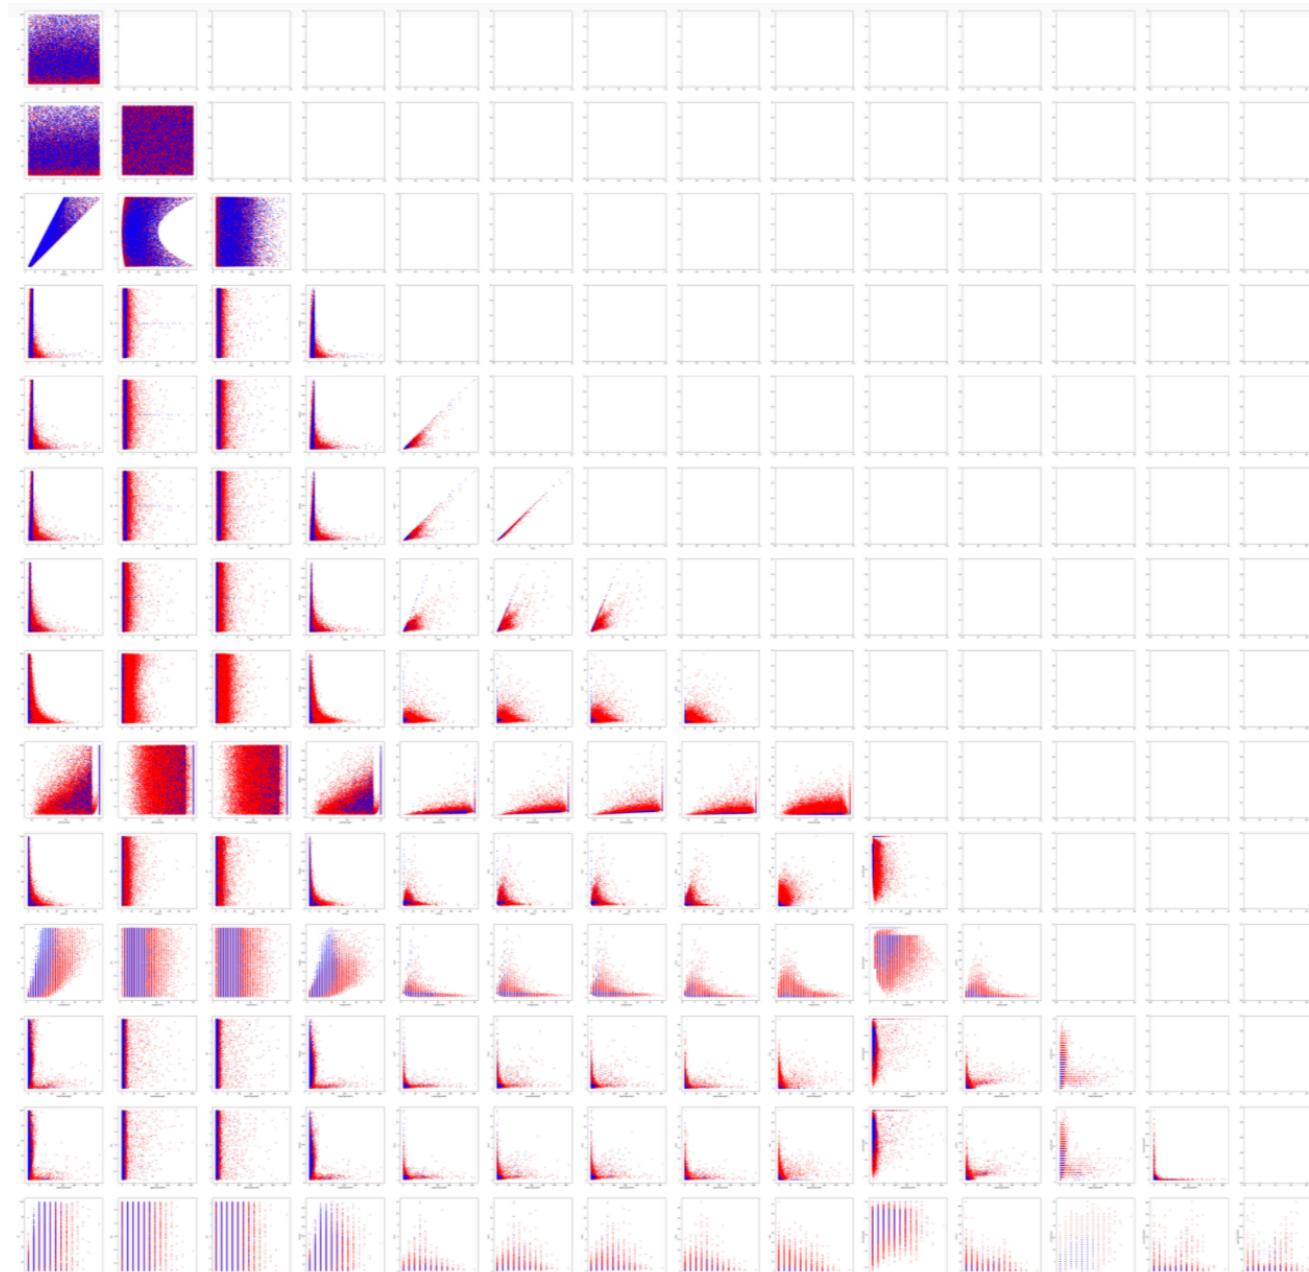


Dataset Visualization



- Start with pattern of detector hits
- Two main tasks:
 - Associate each incident particle with a collection of hits

Designing ML



How can I predict if a point is red or blue given x_1 and x_2 and x_3 and ...?
(Lets use the notebook)

Batch Norm & Dropout

- In addition we often employ these two to improve perf

Batch Norm makes output Gaussian-like

Input: Values of x over a mini-batch: $\mathcal{B} = \{x_1 \dots m\}$;
Parameters to be learned: γ, β

Output: $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

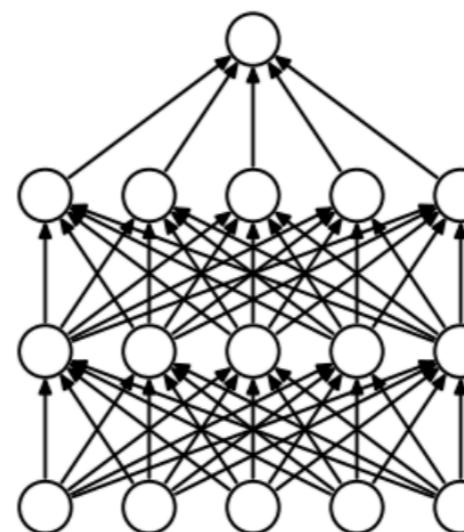
$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{mini-batch variance}$$

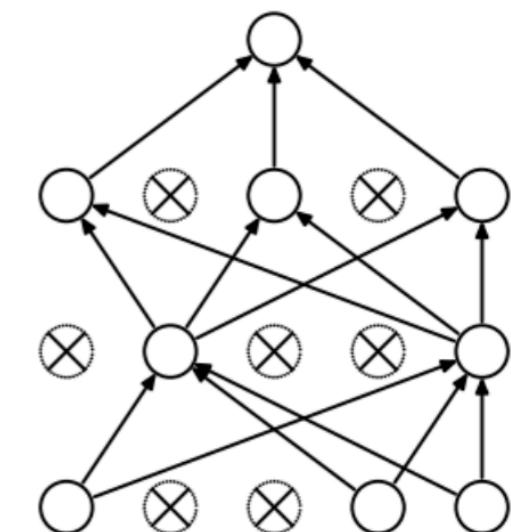
$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{scale and shift}$$

Dropout randomly removes weights



(a) Standard Neural Net



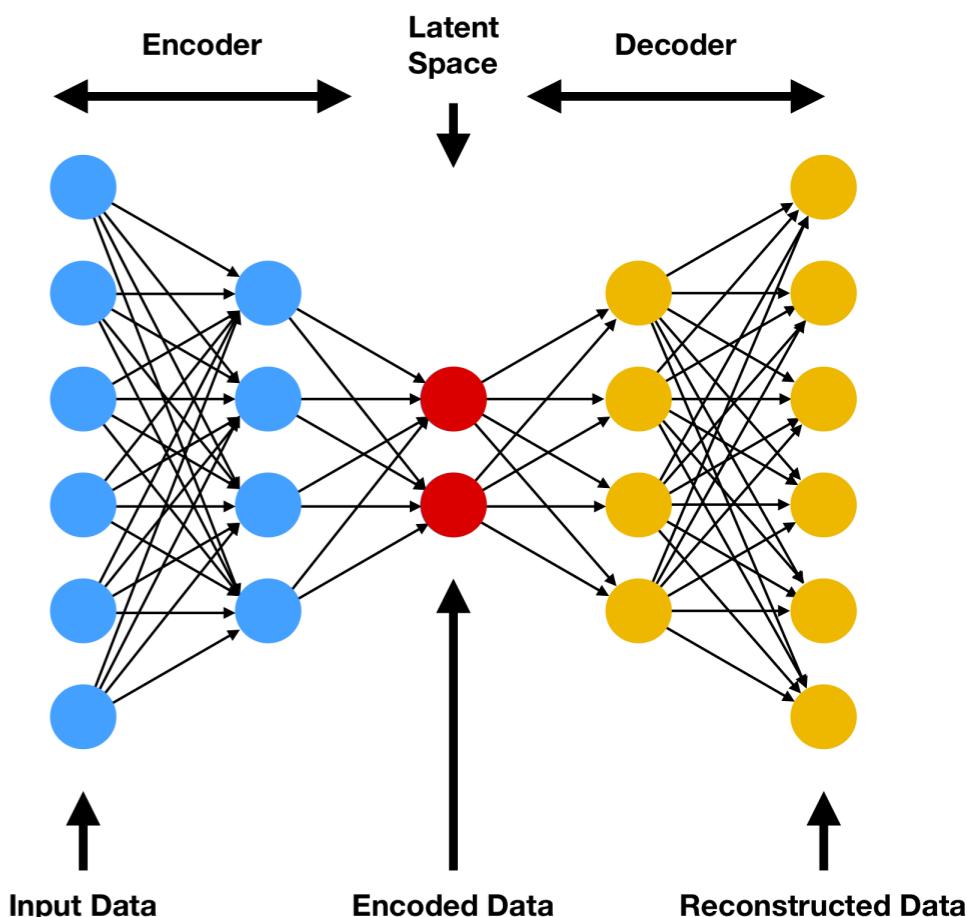
(b) After applying dropout.

- Both Strategies make NNs more robust to deviations

More Complex Architectures

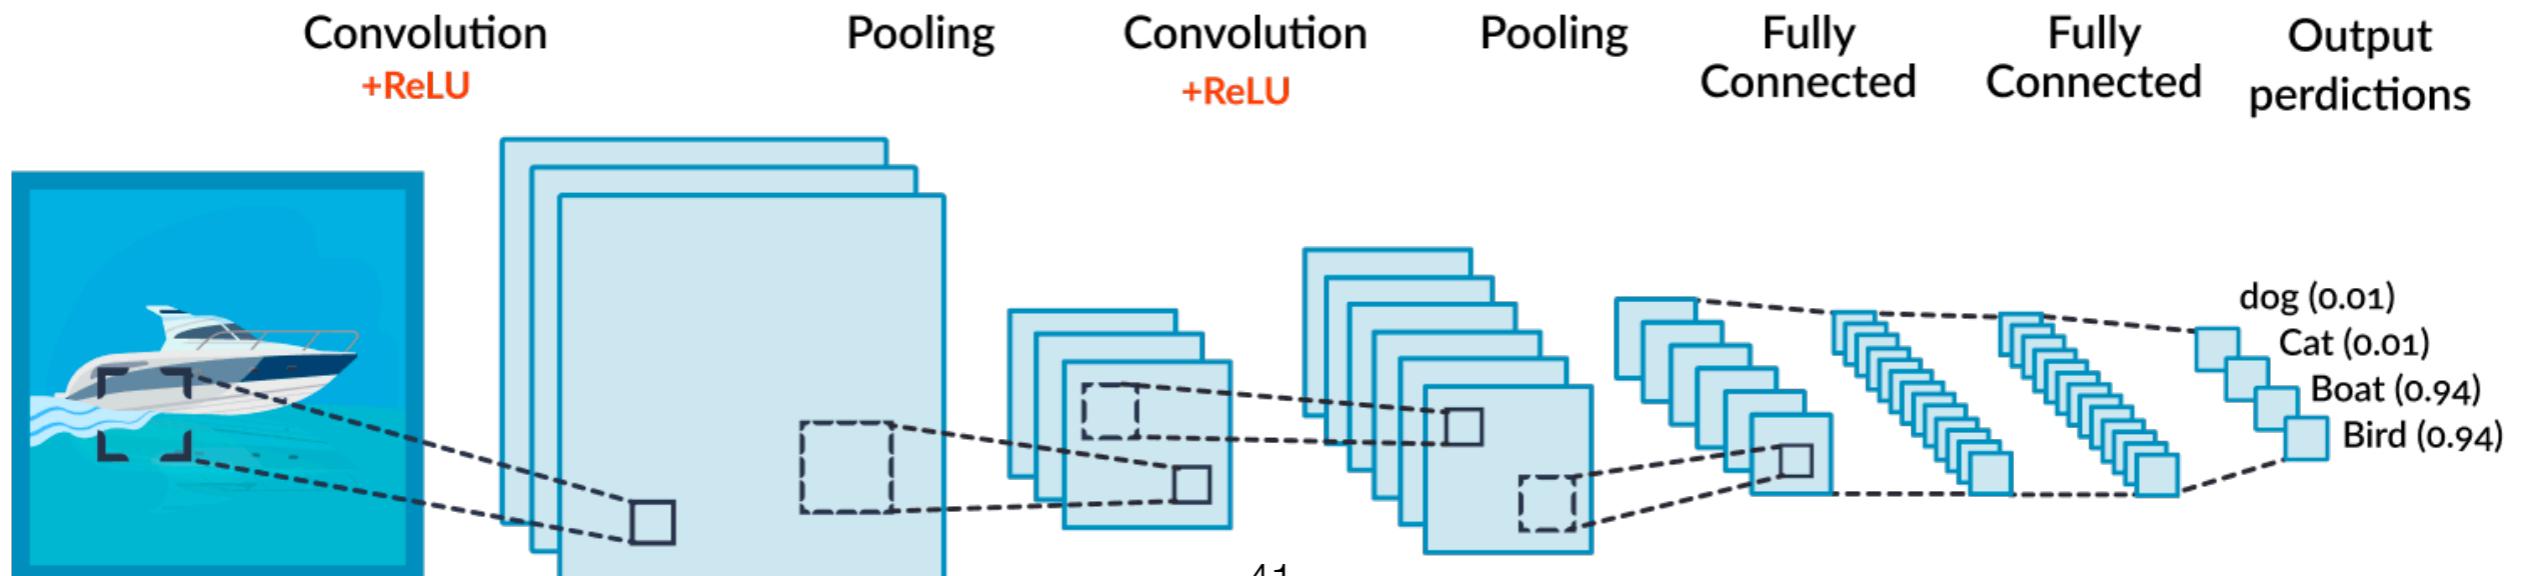
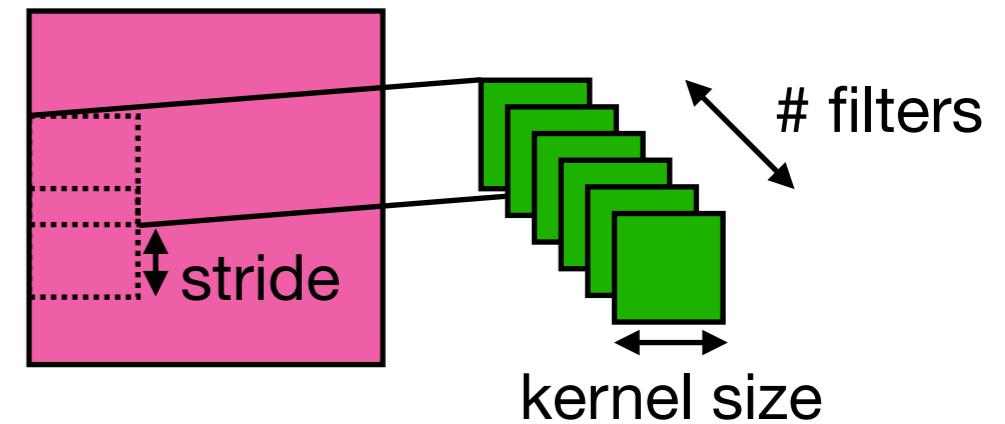
Unsupervised Learning

- What if we don't have/use labels?
- Autoencoder (AE):
 - $\text{Loss} = \text{output} - \text{input}$
 - Latent space can be used for clustering
 - If one class of data is used to train, different classes may not reconstruct well (anomaly detection)



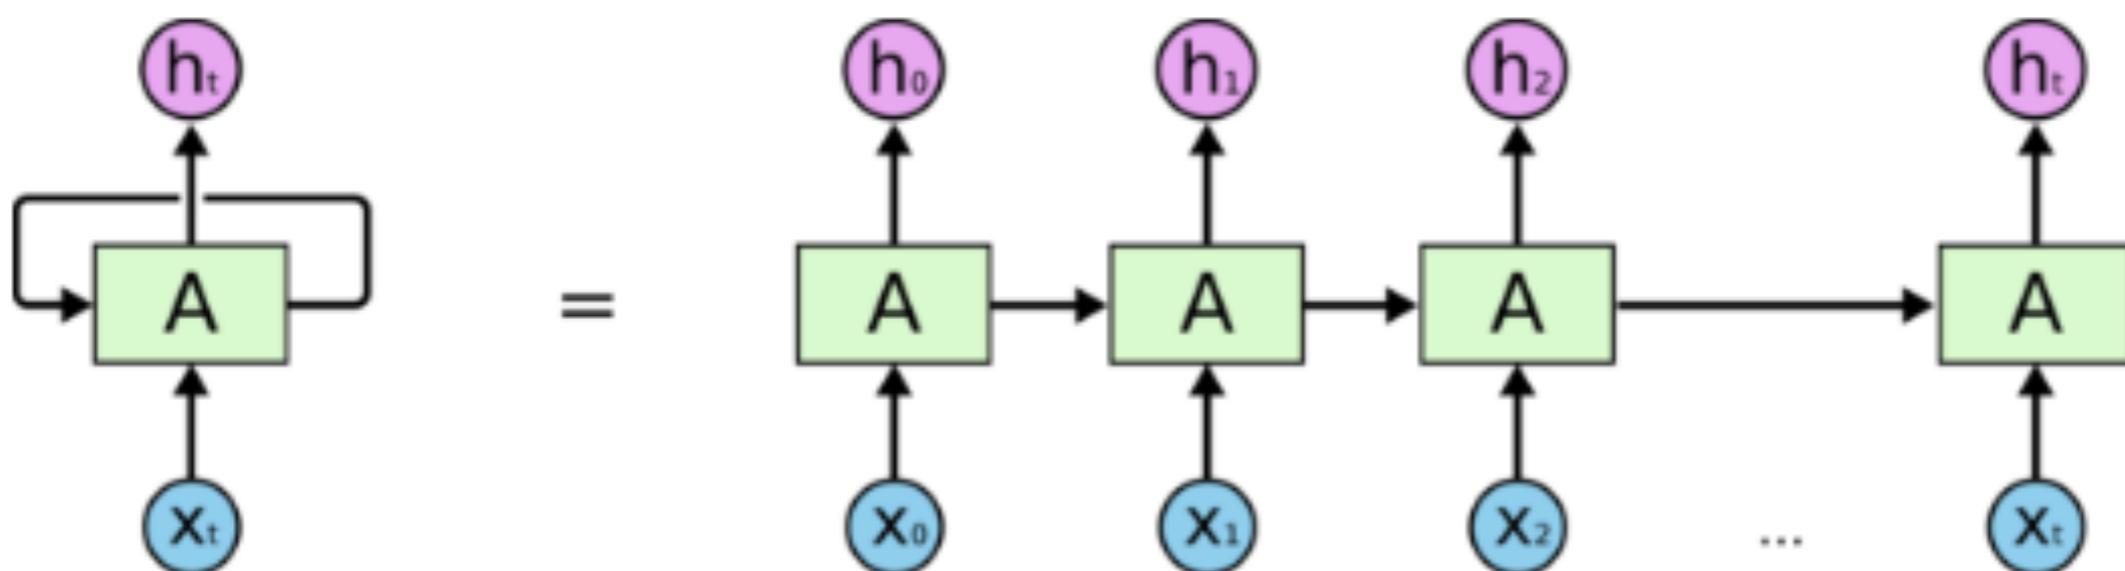
Convolutional Neural Network (CNN)

- Used extensively for images (Conv2D), also useful for 1D and 3D cases
- Small dense network takes a local region as input, scans over whole image
 - *# filters, kernel size, stride*
 - Typically followed by Pooling layer to reduce dimensionality



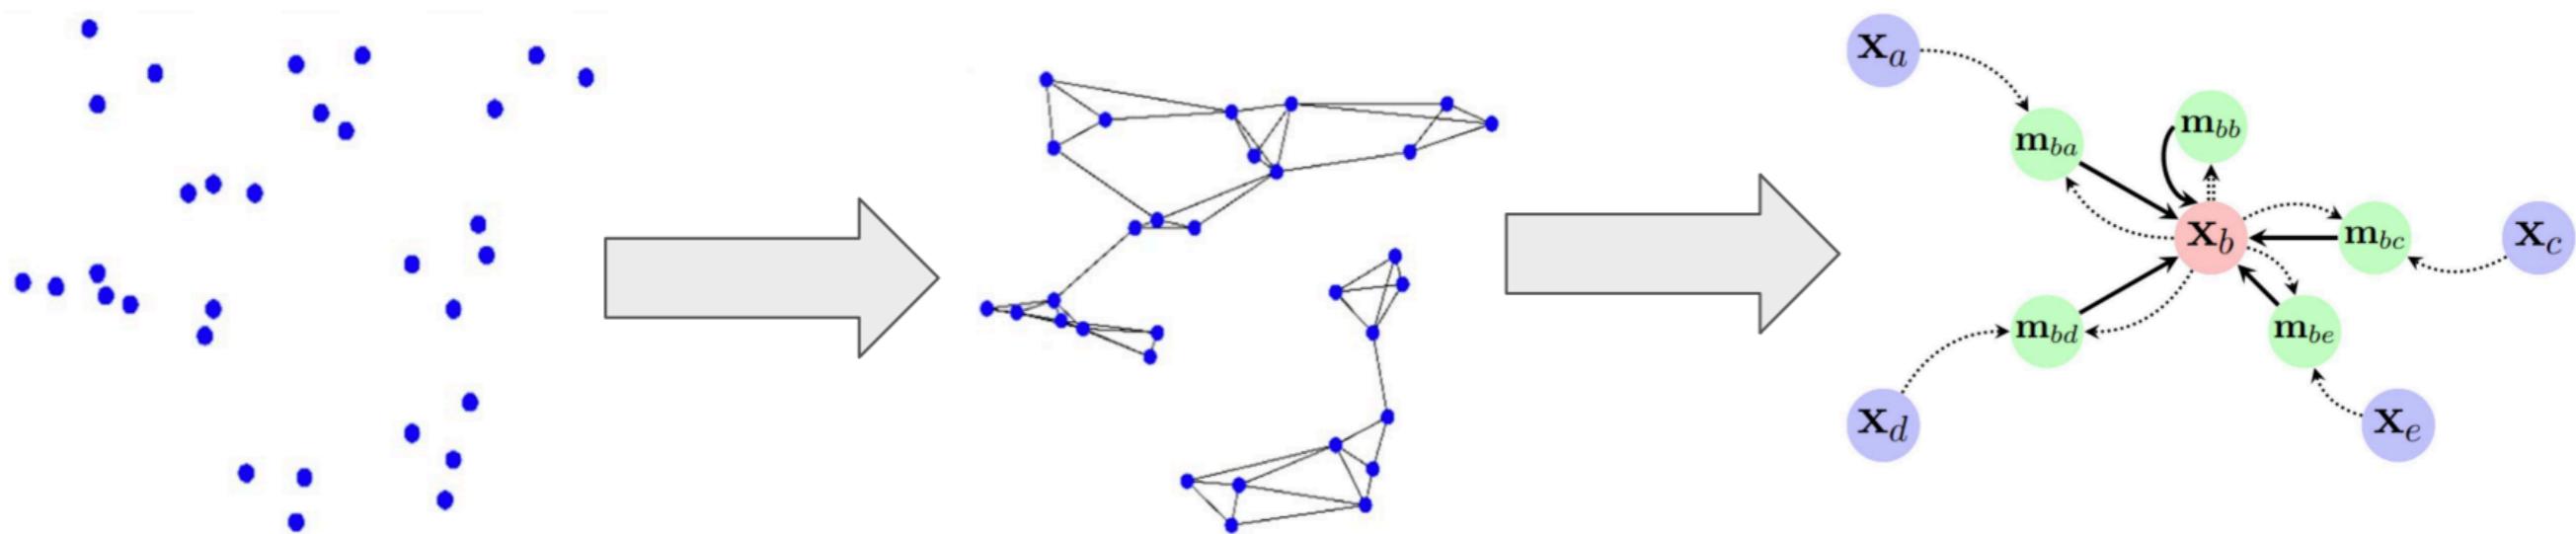
Recurrent Neural Network (RNN)

- Designed for sequential inputs (ex. language)
 - Retain “memory”
- Long short-term memory (LSTM), gated recurrent unit (GRU)



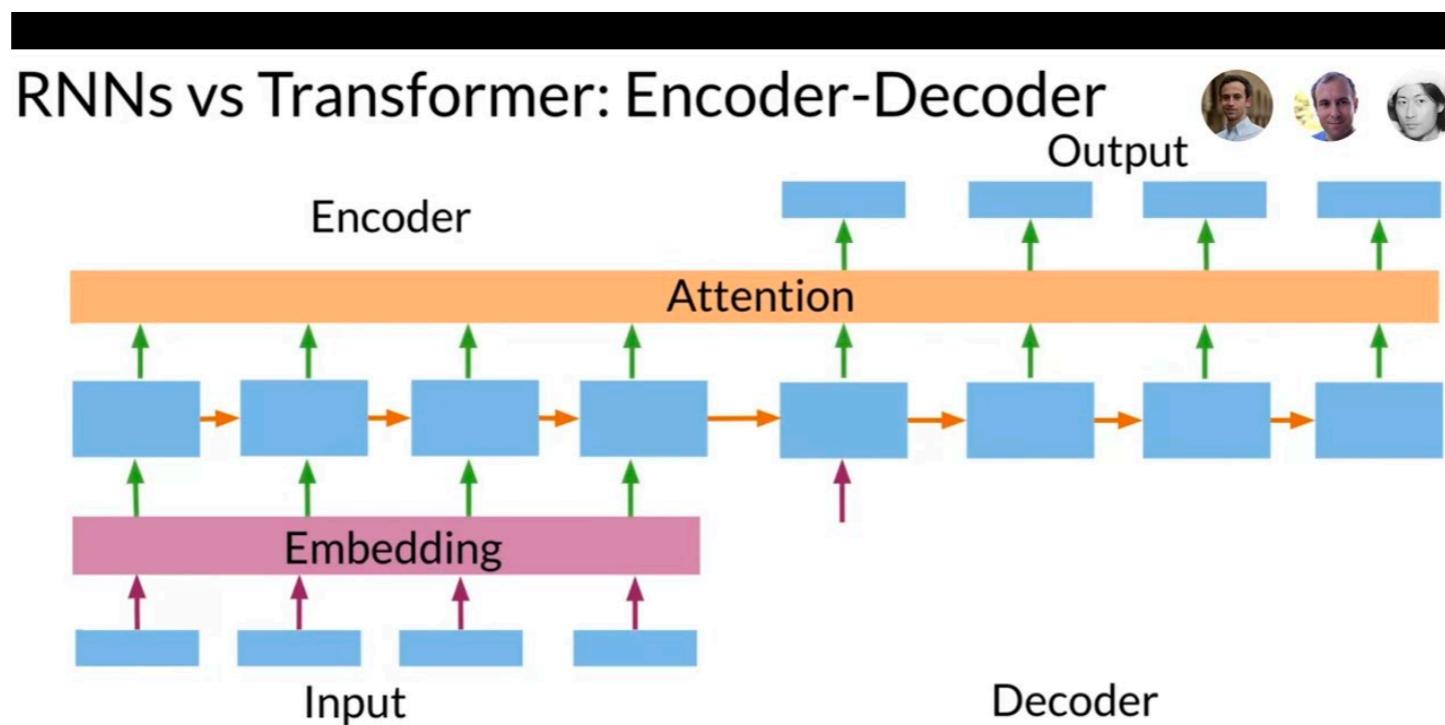
Graph Neural Network (GNN)

- Take in a set of points
 - Construct a graph out of these points
 - Make links between the neighbors
 - Iterate back and forth with the neighbors



Transformers

- Originated from recursive neural networks
 - Reading a whole sentence is better than iteratively
 - Done by applying Attention (effectively a big linear layer)
 - Basically a Graph with a notion of ordering



Next Lecture

- We are going to look at another application of ML
 - Using ML for regressions
 - How can we solve for complex physical scenarios